
ciphart

memory-hard key derivation

with easier measurable security

caveman¹
2021-02-14 08:10:33+00:00

argon2² is mostly nice, but trying to interpret its contribution to the protection against password brute-forcing attacks remains more difficult than it should be. this vagueness is a problem that is not limited to *argon2*, but also shared with every other key derivation function that i've known so far.

when one uses *argon2*, his derived key will surely have superior protection against password brute-forcing attacks, but by how much? to answer this, one would need to survey the industry that manufactures application-specific integrated circuits (asics) to obtain a map between *time* and *money*, in order to get an estimation on how much would it cost the adversary to discover the password in a given time window.

while the approach of surveying the asics industry is not wrong, it is largely subjective, with expensive housekeeping, and practically leads the user to rely on vague foundations to build his security on. this vagueness is not nice, and it would be better if we had an objective measure to quantify the security of our memory-hard key derivation functions.

resolving this vagueness is not a mere luxury to have, but a necessity for maximising survival, because it hinders the process of studying the cost-value of memory-hard key derivation functions, which, effectively, increases the risk of having a false sense of security.

so i propose *ciphart* — a memory-hard key derivation function with a security contribution that is measured in a unit that i call *caveman's entropy bits*. this unit is measured objectively and is guaranteed to be true irrespective of whatever alien technology that the adversary might have.

`libciphart`³ is a library that implements *ciphart* very closely to this paper, without much fluff. this should make integrating *ciphart* into other systems more convenient.

`ciphart`⁴ is an application for encrypting and decrypting files that makes use of `libciphart`. this application is intended for use by end-users or scripts, henceforth it has some fluff to treat mankind with dignity.

paper's layout

1 shannon's entropy

¹mail: toraboracaveman [at] protonmail [dot] com

²<https://github.com/P-H-C/phc-winner-argon2>

³<https://github.com/Al-Caveman/libciphart>

⁴<https://github.com/Al-Caveman/ciphart>

2	caveman's entropy	2
2.1	recursive <code>hash</code>	2
2.2	memory-hard <code>hash</code>	2
2.3	implications	2
3	ciphart	3
3.1	parameters	3
3.2	internal variables	3
3.3	output	3
3.4	steps	3
4	parallelism	4
5	memory-hardness	4
6	security interpretation	4
6.1	key brute-forcing	4
6.2	normal password brute-forcing	5
6.3	with <i>argon2</i>	5
6.4	with <i>ciphart</i>	6
7	summary	6

1 shannon's entropy

we've got password p with $H(p)$ many shannon's entropy bits worth of information in it. so what does this mean?

fundamentally, it means that, on average, we'd need to ask $H(p)$ many perfect binary questions⁵ in order to fully resolve all ambiguities about p ; i.e. to fully get every bit of p .

but people use it to do less orthodox things, such as quantifying the amount of security p has against, say, brute-forcing attacks.

say that we've got a 8V bit key $k \leftarrow \text{hash}(p\|s, 8V)$, derived from password p , where s is a salt. say that the attacker has s and k but wants to figure out p . in this case, he will need to brute-force the password space in order to find p that gives k . his cost is:

$$2^{H(p)} \left(\text{cost}(\text{hash}) + \text{cost}(\text{if } \hat{k} = k) \right) \quad (1)$$

definition 1. *the security of a system is the cost of the cheapest method that can break it.*

one way to estimate `cost` is to survey the asics industry. by surveying the asics industry to get an idea how much money it costs to get a given key, or password, space brute-forced within a target time frame⁶. this has an expensive housekeeping and is usually not possible to get any guarantees as we don't know about state-of-art manufacturing secrets that adversaries may have. the *scrypt* paper has an example of such attempt.

⁵one which, if answered, and on average, gets the search space reduced in half.

⁶see the *scrypt* paper for an example.

another way is to ignore anything that has no cryptographic guarantee. so, in (1), cryptography guarantees⁷ that $2^{H(p)}$ many **hash** calls are performed and that many equality tests. the **hash** call needs to be done once, so let's give it a unit of time 1. the equality test also needs to be called once, but since it's so cheap it's easier to just assume that its cost is free. this way (1) becomes just:

$$2^{H(p)}(1 + 0) = 2^{H(p)} \quad (2)$$

i think this is why people use password entropy as a measure of its security. not because it is the quantity of security, but rather because it's the quantity of *simplified* security. further, for convenience, it seems that people report it in the \log_2 scale. i.e. $\log_2 2^{H(p)} = H(p)$.

2 caveman's entropy

2.1 recursive hash

if the **hash** function is replaced by an N -deep recursion over **hash**, like:

$$\begin{aligned} & \text{rhash}(p||s, 8V, N) \\ &= \text{hash}(\text{hash}(\dots \text{hash}(p||s, 8V), \dots, 8V), 8V) \end{aligned}$$

then, if **hash** is not broken, (1) becomes:

$$2^{H(p)} \left(N \text{cost}(\text{hash}) + \text{cost}(\text{if } \hat{k} = k) \right) \quad (3)$$

and (2) becomes:

$$\begin{aligned} 2^{H(p)}(N + 0) &= N 2^{H(p)} \\ &= 2^{H(p) + \log_2 N} \end{aligned} \quad (4)$$

at this point, thanks to cryptographic guarantees, there is absolutely no security distinction between a password with shannon's $H(p) + \log_2 N$ entropy bits, and a password with just $H(p)$ entropy bits that made use of the N -deep recursive calls of **hash**.

shannon's entropy of p remains $H(p)$, but thanks to the recursive calls of **hash**, that password will be as expensive as another password \hat{p} , such that $H(\hat{p}) = H(p) + \log_2 N$.

i think it will be simpler if we introduce the function-dependent caveman's entropy C as a measure. it goes like this:

$$C(p, \text{hash}(\dots)) = H(p) \quad (5)$$

$$C(\hat{p}, \text{hash}(\dots)) = H(p) + \log_2 N \quad (6)$$

$$\begin{aligned} C(p, \text{rhash}(\dots, N)) &= H(p) + \log_2 N \\ &= H(\hat{p}) \end{aligned} \quad (7)$$

security-wise, there is no distinction between the more complex password \hat{p} , and the simpler password p that used **rhash**(..., N). so i really think we need to measure password security in C instead of H .

⁷statistically by confidence earned through peer review and attempts to break encryption algorithms.

2.2 memory-hard hash

let **mhash** be like **rhash**, except that it also requires M many memory bites such that, as available memory is linearly reduced from M , penalty in cpu time grows exponentially. let M be requested memory, \hat{M} be available memory, and $e(M - \hat{M})$ be the exponential penalty value for reduction in memory, where $e(0) = 1$.

$$\text{mhash}(p||s, N, M) = \text{rhash}(p||s, N)^{e(\hat{M} - M)} \quad (8)$$

if the **hash** function is replaced by M bytes memory-hardened N -deep recursion hash function **mhash**, (1) becomes:

$$2^{H(p)} \left(N^{e(M - \hat{M})} \text{cost}(\text{hash}) + \text{cost}(\text{if } \hat{k} = k) \right) \quad (9)$$

(2) becomes:

$$\begin{aligned} 2^{H(p)}(N^{e(M - \hat{M})} + 0) &= N^{e(M - \hat{M})} 2^{H(p)} \\ &= 2^{H(p) + \log_2 N^{e(M - \hat{M})}} \\ &= 2^{H(p) + e(M - \hat{M}) \log_2 N} \end{aligned} \quad (10)$$

and caveman's entropy becomes:

$$C(p, \text{mhash}(\dots, N, M)) = H(p) + e(M - \hat{M}) \log_2 N \quad (11)$$

2.3 implications

let p be a password with $H(p)$ shannon's entropy bits. let \hat{p} be a more complex password with $H(p) + e(M - \hat{M}) \log_2 N$ shannon's entropy bits, where M , \hat{M} and N are all positive numbers.

then caveman's entropy says that the following keys are information theoretically indistinguishable for as long as only p and \hat{p} remain unknown (everything else is known, such as the distribution from which p and \hat{p} was sampled), and for as long as **hash** is not broken:

- $k \leftarrow \text{mhash}(p||s, N, M)$
- $\hat{k} \leftarrow \text{hash}(\hat{p}||s)$

in other words:

$$C(p, \text{mhash}(\dots, N, M)) = H(\hat{p}) \quad (12)$$

since the assumption that passwords are kept away from the adversary is fundamental in a symmetric encryption context, i think it makes sense that we measure our security with memory-hard key derivation functions using the caveman's entropy C .

from a security point of view, it will feel absolutely identical to as if the password got injected with extra shannon's entropy bits. no one can tell the difference for as long as the fundamental assumption of hiding passwords is honoured, as well as the hashing function **hash** is not broken.

in other words, we can say, if password p is unknown, and **hash** is not broken, then we have injected into p extra shannon's entropy bits. this lie will be only discovered

after p is revealed — call it caveman’s cat thought experiment.

if you think that it is impossible for this *lie* to be *truth* under the secrecy of p , then i’ve done an even better job: proving that cryptographically secure hashing functions do not exist. likewise, same can be trivially extended to: cryptographically symmetric ciphers do not exist.

so you have to pick either one of two options:

1. either accept that the lie is truth. i.e. we injected shannon’s entropy bits into p , for as long as only p is not revealed.
2. or, accept that cryptographically-secure hashing and symmetric-encryption functions do not exist.

theorem 1 (the perfect lie). *when p is secret and **hash** is not broken, then shannon’s entropy H of the derived key equals caveman’s entropy C .*

i call theorem 1 the *perfect lie theorem* in a sense that a perfect lie is indistinguishable from truth.

3 ciphart

3.1 parameters

enc	encryption function.
p	password.
s	salt.
M	total memory in bytes.
L	number of memory lanes for concurrency.
T	number of tasks per lane segment.
B	minimum quantity of increased protection against password brute-forcing attacks in the unit of <i>caveman’s entropy bits</i> .
K	output key size in bytes.

3.2 internal variables

hash	hashing function.
C	$\leftarrow \begin{cases} 64 \text{ bytes} & \text{if } \mathbf{enc} \text{ is } xchacha20 \\ 16 \text{ bytes} & \text{if } \mathbf{enc} \text{ is } aes \\ \dots \end{cases}$ this to reflect the block size of the encryption algorithm that implements enc .
V	$\leftarrow \begin{cases} 32 \text{ bytes} & \text{if } \mathbf{enc} \text{ is } xchacha20 \\ 16 \text{ bytes} & \text{if } \mathbf{enc} \text{ is } aes-128 \\ 32 \text{ bytes} & \text{if } \mathbf{enc} \text{ is } aes-256 \\ \dots \end{cases}$ this is the size of the encryption key that’s used to solve <i>ciphart</i> ’s tasks. this is different than the enc -independent K which is possibly used by other encryption algorithms in later stages ⁸ .
\hat{T}	$\leftarrow \max(\lceil VC^{-1} \rceil, T)$. this is to ensure that we have enough encrypted bytes for new keys.
\hat{T}	$\leftarrow \hat{T} - (\hat{T} \bmod 2) + 2$. this is to ensure that there is an even number of tasks in a segment. why? because we need a buffer for storing the clear-text and another for storing the output cipher-text.
\hat{M}	$\leftarrow M - (M \bmod C\hat{T}L) + C\hat{T}L$. this is to ensure that it is in multiples of $C\hat{T}L$. why? so that all segments are of equal lengths in order to simplify <i>ciphart</i> ’s logic. e.g. it wouldn’t be nice if the last segments were of unequal sizes.
G	$\leftarrow \hat{M}C^{-1}\hat{T}^{-1}L^{-1}$. total number of segments per lane.
N	$\leftarrow 0$. actual number of times enc is called, where $\hat{N} \geq 2^B$.
m_i	C -bytes memory for i^{th} task in the \hat{M} -bytes pad.
n_l	$\leftarrow lG\hat{T}$. nonce variable for l^{th} lane with at least 64 bits.
f	$\leftarrow 0$. a flag indicating whether the \hat{M} -bytes pad is filled.
v	$\leftarrow *hash(p \parallel s, V)$. a pointer to the first byte where V -bytes key is stored.

3.3 output

k	K -bytes key.
\hat{B}	actual quantity of increased security against password brute-forcing attacks in the unit of <i>caveman’s entropy bits</i> , where $\hat{B} \geq B$.

3.4 steps

steps of *ciphart* is shown in algorithm 1. this corresponds to *argon2d*. adding a *ciphart-i* variant is a trivial matter, i just didn’t do it yet because my threat model currently doesn’t benefit from a password independent variant.

⁸at the expense of losing the meaning of *caveman’s entropy bits*.

algorithm 1: ciphart

```
1 while 1 do
2   for  $g = 0, 1, \dots, G - 1$  do
3     for  $l = 0, 1, \dots, L - 1$  do
4       for  $t = 0, 1, \dots, T - 1$  do
5          $i \leftarrow gLT + lT + t$ ;
6         if  $t < T - 1$  then
7            $j \leftarrow i + 1$ ;
8         else if  $t = T - 1$  then
9            $j \leftarrow i - T + 1$ ;
10         $m_j \leftarrow \text{enc}(m_i, n_l, v)$ ;
11         $n_l \leftarrow n_l + 1$ ;
12        if  $f = 0$  then
13           $v \leftarrow m_j \bmod (gLTC + tC - V)$ ;
14          if  $v \geq gLTC - V$  then
15             $v \leftarrow v + LTC$ ;
16        else
17           $v \leftarrow m_j \bmod (\hat{M} - LTC + tC - V)$ ;
18          if  $v \geq gLTC + tC - V$  then
19             $v \leftarrow v + LTC$ ;
20          else if  $v \geq gLTC - V$  then
21             $v \leftarrow v + LTC$ ;
22       $N \leftarrow N + LT$ ;
23      if  $N \geq 2^B$  then
24         $g_{\text{last}} \leftarrow g$ ;
25        go to line 27;
26   $f \leftarrow 1$ ;
27  $i \leftarrow g_{\text{last}}LT$ ;
28  $k \leftarrow \text{hash}(m_{i+0T} \| m_{i+1T} \| \dots \| m_{i+(L-1)T}, K)$ ;
29  $\hat{B} \leftarrow \log_2 N$ ;
30 return  $k, \hat{B}$ 
```

4 parallelism

since iterations of the loop in line 3 in algorithm 1 are fully independent of one other, they can quite happily utilise L cpu cores, specially when segment sizes, T , are larger.

5 memory-hardness

Proof. algorithm 1 is just a variation of *argon2d*, except that it uses an encryption function, **enc**, instead of a hashing function. so if *argon2d* is memory-hard, then so is *ciphart*. \square

6 security interpretation

note 1. *i assume that the decryption part of the encryption algorithm **enc** costs the same as the encryption. this is true for algorithms such as xchacha20. and in cases where it is not true, such as with aes, ciphart can simply*

encrypt using the decryption function. this way we guarantee that the cost are identical between ciphart's encryption, and the cipher-text decryption that the adversary does to test a given key.

let's say that we used block encryption function **enc** and a key $v \leftarrow \text{hash}(p \| s, V)$ to encrypt some clear-text into a sequence of C -byte cipher-text blocks m_0, m_1, \dots . let's say that the adversary got those m_0, m_1, \dots .

adversary's goal is to decrypt those cipher-text blocks back into the original clear-text. so what options does he have?

6.1 key brute-forcing

brute-force the V -bytes key space. in order to get a probability of 1 of finding the key v , the adversary would need to evaluate 2^{8V} many keys⁹.

this works by having the adversary repeatedly decrypting m_0 with **enc**, each time using a new key among

- $\hat{v}_0 \leftarrow 0x00 \dots 0$,
- $\hat{v}_1 \leftarrow 0x00 \dots 1$,
- \vdots
- $\hat{v}_{2^{8V}-1} \leftarrow 0xff \dots f$,

until the adversary finds a key that manages to decrypt m_0 .

the adversary could be extremely lucky and have v_0 manage to decrypt m_0 , hence needing to call **enc** only once.

or he might be extremely unlucky and need to keep trying until $v_{2^{8V}-1}$ manages to do it, hence needing to call **enc** for 2^{8V} many times.

usually it's sometime in between. asymptotically n on average, the adversary would need to call **enc** for $2^{8V}/2$ many times.

but in order to guarantee finding v , the brute-forcing process would need to run 2^{8V} many evaluations, hence calling **enc** for 2^{8V} many times.

that said, the adversary would be fossilised long before his application completes. e.g. since $8V = 256$ is common for ciphers nowadays, on average while considering the lucky and the unlucky cases, it would take my laptop 4.28×10^{58} centuries to just increment a counter for 2^{256} many times. so if the adversary's best hope is to brute-force keys, our system has reached maximum security.

security interpretation 1. *your protection against key brute-forcing attacks with key brute-forcing is $2^{8V} \text{cost}(\text{enc})$, where $\text{cost}(\text{enc})$ is the cost of executing **enc** a single time.*

*this is usually called 8V entropy bits. but for the purpose of helping later sections, i think it's better to call it 8V entropy bits from the viewing angle of **enc**, or, for short:*

8V HENC

⁹assuming that each byte is 8 bits.

definition 2. HENC is entropy bits from the viewing angle of **enc**.

HENC is specific only to **enc**'s algorithm, so must hold with whatever alien technology that the adversary may have, for as long as **enc**, as an algorithm, has no cryptanalysis. if there is any cryptanalysis, we'll have to subtract bits from $8V$, e.g. $8V - z$ HENC, where z is number of reduced bits due to cryptanalysis.

6.2 normal password brute-forcing

brute-force the $H(p)$ bits password space. where $H(p)$ is the amount of entropy bits in p . in order to get a probability of 1 of finding the password p , the adversary would need to evaluate $2^{H(p)}$ many keys¹⁰.

this works by having the adversary repeatedly decrypting m_0 with **enc**, each time using a new key among:

- $\hat{v}_0 \leftarrow \text{hash}(\hat{p}_0 \| s, V)$,
- $\hat{v}_1 \leftarrow \text{hash}(\hat{p}_1 \| s, V)$,
- \vdots
- $\hat{v}_{2^{H(p)}-1} \leftarrow \text{hash}(\hat{p}_{2^{H(p)}-1} \| s, V)$,

until the adversary finds a key that manages to decrypt m_0 .

security interpretation 2. your protection against password brute-forcing attacks with normal hashed passwords is $2^{H(p)} \text{cost}(\text{hash}) + 2^{H(p)} \text{cost}(\text{enc})$. the latter **enc** calls are due to trying to decrypt m_0 at every attempt.

this is usually called $H(p)$ entropy bits. but for the purpose of this paper, i think it's better to be more specific. from the viewing angle of **enc**, we get the entropy:

$$H(p) \text{ HENC}$$

and from the viewing angle of **hash**, we get the entropy:

$$H(p) \text{ HHASH}$$

this may seem silly as it is too obvious, but i think it helps me to communicate my thoughts in later sections.

definition 3. HHASH is entropy bits from the viewing angle of **hash**.

6.3 with argon2

adversary evaluates keys from:

- $\hat{v}_0 \leftarrow \text{argon2}(\hat{p}_0, N, M, \dots)$,
- $\hat{v}_1 \leftarrow \text{argon2}(\hat{p}_1, N, M, \dots)$,
- \vdots

- $\hat{v}_{2^{H(p)}-1} \leftarrow \text{argon2}(\hat{p}_{2^{H(p)}-1}, N, M, \dots)$,

for every *argon2* call, **hash** is called for N many times if there is M bytes of memory. so, from the viewing angle of **hash**, we get:

- $\hat{v}_0 \leftarrow \text{hash}(\hat{p}_0 \| s_0, V)$,
- $\hat{v}_1 \leftarrow \text{hash}(\hat{p}_1 \| s_1, V)$,
- \vdots
- $\hat{v}_{N2^{H(p)}-1} \leftarrow \text{hash}(\hat{p}_{N2^{H(p)}-1} \| s_{N2^{H(p)}-1}, V)$,

if an adversary lacks M bytes, he can still compute *argon2*(p, N, M, \dots), but at the expense of needing exponentially more cpu time as his memory is linearly reduced.

security interpretation 3. your protection against password brute-forcing attacks under the *argon2* protection is:

$$\begin{aligned} & N2^{H(p)} \text{cost}(\text{hash}) + 2^{H(p)} \text{cost}(\text{enc}) \\ &= 2^{H(p) + \log_2 N} \text{cost}(\text{hash}) + 2^{H(p)} \text{cost}(\text{enc}) \end{aligned}$$

from the viewing angle of **enc** we get the entropy:

$$H(p) \text{ HENC}$$

while from the viewing angle of **hash** we get the entropy:

$$H(p) + \log_2 N \text{ HHASH}$$

so which one to pick? i think people so far just pick $H(p) \text{ HENC}$ to reflect password's entropy, and seem to not pick $H(p) + \log_2 N \text{ HHASH}$ as they don't seem to consider it entropy. but i have two disagreements with people:

- i think not accepting that $H(p) + \log_2 N \text{ HHASH}$ is entropy is needlessly limiting. because i think $H(p) + \log_2 N \text{ HHASH}$ is entropy as much as $H(p) \text{ HENC}$ is entropy; it's just that they are measured from different viewing angles: former is measured from the **hash** viewing angle, while the latter is measured from the **enc** viewing angle.
- i don't see any reason why any of them is more true than the other. i think that both of them are entropies, but of different units.
- why do people only pick either one of them? it's technically false my view. in my view truth is: we're just dealing with two entropies measured in different units. so i think truth is that we have the following number of entropy bits:

$$\begin{aligned} & H(p) \text{ HENC} \\ & + H(p) + \log_2 N \text{ HHASH} \end{aligned}$$

which obviously looks a bit ugly, since we cannot sum them due to the terms having different units, which also gives our brain a hard time to get a feeling of what that even means.

¹⁰the adversary does not know p , obviously, but he knows the process that the user used to generate p , henceforth he knows $H(p)$.

so what's the solution here to this ugliness? should we ignore $H(p)$ HENC $+H(p) + \log_2 N$ HHASH as an entropy measure that quantifies the security of our protection against password brute-forcing, as people currently do, and measure it only in terms of the computational cost by surveying the industry of asics in order to find a map between time and money?

my answer to the questions above is:

- no. the right approach is to just admit that *argon2*'s approach is dragging us into the situation where we end up with two entropies measured in different units.
- *argon2*'s security contribution is measurable as entropy, except that it is ugly since it is made of two entropies in distinct units. if we ignore this, we won't solve the problem, but end up stashing the dirt under the carpet.
- of course, we are always free to also survey the industry of asics to derive time-money maps, but this doesn't have to be our only approach to quantify our security gain.

6.4 with *ciphart*

adversary evaluates keys from:

- $\hat{v}_0 \leftarrow \text{ciphart}(\hat{p}_0, B, M, \dots),$
- $\hat{v}_1 \leftarrow \text{ciphart}(\hat{p}_1, B, M, \dots),$
- \vdots
- $\hat{v}_{2^{H(p)}-1} \leftarrow \text{ciphart}(\hat{p}_{2^{H(p)}-1}, B, M, \dots),$

mostly similar to *argon2*. differences related to this section is that *ciphart* calls *enc* instead of *hash*, and specifies B instead of N , where $B \approx \log_2 N$. similar exponential time penalty applies with memory less than M .

security interpretation 4. *your protection against password brute-forcing attacks under the ciphart protection approach is:*

$$\begin{aligned}
& \left(2^{H(p)} + 2^{\hat{B}} \right) \text{cost}(\text{enc}) + 2^{H(p)} \text{cost}(\text{enc}) \\
&= \left(2^{H(p)} + 2^{\hat{B}} + 2^{H(p)} \right) \text{cost}(\text{enc}) \\
&= \left(2 \times 2^{H(p)} + 2^{\hat{B}} \right) \text{cost}(\text{enc}) \\
&= \left(2^{H(p)+\log_2 2} + 2^{\hat{B}} \right) \text{cost}(\text{enc}) \\
&= \left(2^{H(p)+1} + 2^{\hat{B}} \right) \text{cost}(\text{enc})
\end{aligned}$$

from the viewing angle of *enc* we get the entropy:

$$H(p) + 1 + \hat{B} \text{ HENC}$$

and there is no other viewing angle than *enc* since only *enc* is used! as a result our brain can easily interpret it.

plus, if we wish to study the industry of asics to obtain time-money maps, our job will be much easier as we can simply look at the cost of asics that are already implemented for *enc*¹¹.

7 summary

¹¹e.g. if *enc* is a popular algorithm, such as aes-256, then we can get more specific data from manufacturers, ultimately giving us a more accurate time-money maps. but when *enc* is not popular, we may need to do rougher calculations based on the expected asics area as done in the *script* paper.