

ciphart

faster memory-*harder* key derivation with easier security interpretation

caveman¹

2021-03-05 23:09:08+00:00

synopsis—*argon2*² is a fast and simple memory-hard key derivation function. compared to *scrypt*³, *argon2* is better, specially for its simplicity. but i claim that *argon2* is not fast enough, not memory-hard enough, and its contribution to our security is not simple enough to understand.

henceforth, i propose *ciphart*, which is:

- easier — because its security contribution is measured in the unit of shannon’s entropy. i.e. when *ciphart* derives a key for you, it tells you that it has *injected* a specific guaranteed quantity of shannon’s entropy bits into your derived key. this is possible thanks to my invention, the “entropy injection” theorem.

this offers a great help as it gives us yet another much simpler approach to quantify our security gain as opposed to being limited to surveying the industry of application-specific integrated circuits as done in the *scrypt* paper.

- harder — because it can require crazy-large amounts of memory, beyond our random-access memory, thanks to it being able to use the hard-disk as well. this is possible thanks to my discovery “cacheable keys”.

this is optional, but i extremely like it as it effectively gives me much more security while eventually becoming much faster as well, and the adversary cannot get my cache even if he steals my hard-disks.

- faster — because it does not abuse hashing functions. it uses hashing functions when using them is more suited, and uses symmetric block encryption functions when using them is more suited. this is thanks to my “hashing is only for compression” law.

argon2 incorrectly limits itself to only use a hashing function. at the surface it may appear simpler, but it is actually more complex as it ends up re-inventing what resembles a symmetric block encryption function off the hashing function, except for being slower and with potential entropy loss.

*libciphart*⁴ is a library that implements *ciphart* very closely to this paper, without much fluff. this should make integrating *ciphart* into other systems more convenient.

*ciphart*⁵ is an application for encrypting and decrypting files that makes use of *libciphart*. this application is intended for use by end-users or scripts, henceforth it has some fluff to treat mankind with dignity.

¹mail: toraboracaveman [at] protonmail [dot] com.

²<https://github.com/P-H-C/phc-winner-argon2>

³<http://www.tarsnap.com/scrypt/scrypt.pdf>

⁴<https://github.com/Al-Caveman/libciphart>

⁵<https://github.com/Al-Caveman/ciphart>

paper’s layout

1	background	1
1.1	shannon’s entropy	1
1.2	caveman’s entropy	2
1.2.1	recursive hash: rhash	2
1.2.2	memory-hard hash: mhash	2
2	fundamental ideas	2
2.1	“entropy injection” theorem	2
2.1.1	mental warm up	2
2.1.2	the real deal	3
2.2	“cacheable keys” discovery	3
2.2.1	why does it work	3
2.2.2	potential adversary strategies	4
2.3	“hashing is only for compression” law	4
3	ciphart	5
3.1	the algorithm	5
3.1.1	parameters	5
3.1.2	internal variables	5
3.1.3	output	6
3.1.4	steps	6
3.2	noteworthy features	6
3.2.1	parallelism	6
3.2.2	memory-hardness	6
3.2.3	memory-hardness	6
3.3	comparison	7
4	application scenarios	7
4.1	a currently-useful scenario	7
4.2	a later-useful scenario	7
A	donations	8
A.1	bitcoin	8

1 background

1.1 shannon’s entropy

we’ve got password p with $H(p)$ many shannon’s entropy bits worth of information in it. so what does this mean⁶?

fundamentally, it means that, on average, we’d need to ask $H(p)$ many perfect binary questions⁷ in order to fully resolve all ambiguities about p ; i.e. to fully get every bit of p .

but people use it to do less orthodox things, such as quantifying the amount of security p has against, say, brute-forcing attacks.

say that we’ve got a $8V$ bit key $k \leftarrow \text{hash}(p\|s, 8V)$, derived from password p , where s is a salt. say that the attacker has s and k but wants to figure out p . in this case,

⁶say that $H(p)$ was calculated using the base 2 logarithm, \log_2 . all shannon’s entropies in this paper are calculated this way.

⁷one which, if answered, and on average, gets the search space reduced in half.

he will need to brute-force the password space in order to find p that gives k . his cost is:

$$2^{H(p)} \left(\text{cost}(\text{hash}) + \text{cost}(\text{if } \hat{k} = k) \right) \quad (1)$$

definition 1. *the security of a system is the cost of the cheapest method that can break it.*

one way to estimate **cost** is to survey the asics industry. by surveying the asics industry to get an idea how much money it costs to get a given key, or password, space brute-forced within a target time frame⁸. this has an expensive housekeeping and is usually not possible to get any guarantees as we don't know about state-of-art manufacturing secrets that adversaries may have.

another way is to ignore anything that has no cryptographic guarantee. so, in (1), cryptography guarantees⁹ that $2^{H(p)}$ many **hash** calls are performed and that many equality tests. the **hash** call needs to be done once, so let's give it a unit of time 1. the equality test also needs to be called once, but since since it's so cheap it's easier to just assume that its cost is free. this way (1) becomes just:

$$2^{H(p)}(1 + 0) = 2^{H(p)} \quad (2)$$

further, for convenience, it seems that people report it in the \log_2 scale. i.e.:

$$\log_2 2^{H(p)} = H(p) \quad (3)$$

i think this is why people use shannon's entropy of passwords as a measure of their security. not because it is the quantity of security, but rather because its the quantity of *simplified* security.

i like using shannon's entropy as a measure of simplified security quantity, so i'm going to build on it.

1.2 caveman's entropy

1.2.1 recursive hash: rhash

if the **hash** function is replaced by an N -deep recursion over **hash**, like:

$$\begin{aligned} & \text{rhash}(p||s, 8V, N) \\ &= \text{hash}(\text{hash}(\dots \text{hash}(p||s, 8V), \dots, 8V), 8V) \end{aligned}$$

then, if **hash** is not broken, (1) becomes:

$$2^{H(p)} \left(N \text{cost}(\text{hash}) + \text{cost}(\text{if } \hat{k} = k) \right) \quad (4)$$

(2) becomes:

$$\begin{aligned} 2^{H(p)}(N + 0) &= N 2^{H(p)} \\ &= 2^{H(p) + \log_2 N} \end{aligned} \quad (5)$$

⁸see the *scrypt* paper for an example.

⁹statistically by confidence earned through peer review and attempts to break encryption algorithms.

and the \log_2 scaled version becomes:

$$H(p) + \log_2 N \quad (6)$$

obviously (6), which represents the log scaled version of the simplified security of password p when hashed using **mhash**, is no longer equivalent to p 's shannon's entropy as was the case in (3) when the same password p was hashed by using just **hash**.

i think (6) better have a name. i propose to call it caveman's entropy, C , which i define to be function dependent. it goes like this:

$$C(p, \text{hash}(\dots)) = H(p) \quad (7)$$

$$C(p, \text{rhash}(\dots, N)) = H(p) + \log_2 N \quad (8)$$

1.2.2 memory-hard hash: mhash

let **mhash** be like **rhash**, except that it also requires M many memory bytes such that, as available memory is linearly reduced from M , penalty in cpu time grows exponentially. let M be requested memory, A be available memory, and $e(M - A)$ be the exponential penalty value for reduction in memory, where $e(0) = 1$.

$$\begin{aligned} & \text{cost} \left(\text{mhash}(p||s, N, M) \right) \\ &= \text{cost} \left(\text{rhash}(p||s, N) \right)^{e(M-A)} \end{aligned} \quad (9)$$

if **hash** in (1) is replaced by the M -bytes memory-hardened N -deep recursion hash function **mhash**, then (1) becomes:

$$2^{H(p)} \left(N^{e(M-A)} \text{cost}(\text{hash}) + \text{cost}(\text{if } \hat{k} = k) \right) \quad (10)$$

(2) becomes:

$$\begin{aligned} 2^{H(p)}(N^{e(M-A)} + 0) &= N^{e(M-A)} 2^{H(p)} \\ &= 2^{H(p) + \log_2 N^{e(M-A)}} \\ &= 2^{H(p) + e(M-A) \log_2 N} \end{aligned} \quad (11)$$

and caveman's entropy becomes:

$$C(p, \text{mhash}(\dots, N, M)) = H(p) + e(M - A) \log_2 N \quad (12)$$

2 fundamental ideas

2.1 "entropy injection" theorem

2.1.1 mental warm up

say that x is something that all ambiguities about it are resolved after getting, on average, 100 many perfect binary questions, q_0, q_1, \dots, q_{99} , answered. we say that that's its \log_2 -based shannon's entropy, right? i.e. $H(x) = 100$.

what about another thing y , which all ambiguities about it are resolved after getting, on average, 10

many questions, $\hat{q}_0, \hat{q}_1, \dots, \hat{q}_9$, answered, such that each question \hat{q}_i is made of 10 perfect binary questions $q_{i10+0}, q_{i10+1}, \dots, q_{i10+9}$? what's the entropy of y ?

$H(y) = 100$, because simply bundling 10 perfect binary questions together, and calling them \hat{q}_i won't change the fact that the total perfect number of perfect binary questions that need to be answered is still $10 \times 10 = 100$.

so, $H(x) = H(y) = 100$ because question bundling doesn't reduce the information content.

lemma 1. *bundling perfect binary questions, about a thing, into fewer super questions, doesn't reduce thing's information content.*

2.1.2 the real deal

say that p is a password that all ambiguities about it are resolved after getting, on average, $H(p)$ many perfect binary questions, $q_0, q_1, \dots, q_{H(p)-1}$, answered.

case 1 — say that the adversary got p 's hash, $k \leftarrow \text{hash}(p\|s, 8V)$, and that his goal is to find out p by brute-forcing. in this case, the adversary will be considering password candidates $p_0, p_1, \dots, p_{H(p)-1}$, and tests them by asking the following binary questions in order to guarantee finding p :

- q_0 : is $\text{hash}(p_0\|s, 8V) = k$?
- q_1 : is $\text{hash}(p_1\|s, 8V) = k$?
- \vdots
- $q_{H(p)-1}$: is $\text{hash}(p_{H(p)-1}\|s, 8V) = k$?

case 2 — say that mhash was used instead of just hash ; i.e. the adversary got $k \leftarrow \text{rhash}(p\|s, 8V, N)$ instead. in this case, the adversary will seek to answer the following binary questions:

- \hat{q}_0 : is $\text{rhash}(p_0\|s, 8V, N) = k$?
- \hat{q}_1 : is $\text{rhash}(p_1\|s, 8V, N) = k$?
- \vdots
- $\hat{q}_{H(p)-1}$: is $\text{rhash}(p_{H(p)-1}\|s, 8V, N) = k$?

so, you may say that we're still in need to ask, on average, $H(p)$ many perfect binary questions, right? i disagree.

i think what is happening with rhash is that each question $\hat{q}_0, \hat{q}_1, \dots, \hat{q}_{H(p)-1}$ is a *bundled* question, each, containing N many perfect binary questions.

i.e., for any question \hat{q}_i , answering it is equivalent to answering the N -deep recursion over hash :

$$\hat{q}_i: \text{ is } \text{hash}(\text{hash}(\dots \text{hash}(p_i\|s, 8V), \dots, 8V), 8V) = k?$$

which its answer is “yes” if and only if all of the following N many questions are answered by “yes”:

- is $\text{hash}(p_i\|s, 8V) = k_0$?
- is $\text{hash}(k_0\|s, 8V) = k_1$?

- is $\text{hash}(k_1\|s, 8V) = k_2$?
- \vdots
- is $\text{hash}(k_{N-1}\|s, 8V) = k$?

meaning, if any \hat{q}_i question is answered, it necessarily means that we have answered N many q_i -like questions. if answer to \hat{q}_i is “yes”, it means that the answer to each of those N many q_i -like questions is also “yes”. if answer to \hat{q}_i is “no”, it also means that the answer to each of those N many questions is also “no”¹⁰.

in other words, functions rhash and mhash are using cryptography to force the adversary to need to answer more questions beyond $H(p)$. rhash increases questions by a factor of N . mhash increases questions by a factor of $N^{e(M-A)}$.

theorem 1 (entropy injection). *for any password p , and any positive numbers V , N and M :*

$$H(\text{rhash}(p\|s, 8V, N)) = C(p, \text{rhash}(\dots, N))$$

$$H(\text{mhash}(p\|s, 8V, N, M)) = C(p, \text{mhash}(\dots, N, M))$$

2.2 “cacheable keys” discovery

discovery 1 (cacheable keys). *caching keys securely is easily doable, and great security utility exists in doing so for expensively-derived keys.*

2.2.1 why does it work

- when expensively derived keys are cached, only the first key derivation call will be expensive, while subsequent calls will be semi-instantaneous. this effectively allows users to tolerate much more expensive, or secure, key derivation as it only happens during the initial, say, login phase. subsequent use of the extremely expensive key is instantaneous.

so instead of having the user use a somewhat expensive key derivation by waiting say, 3 seconds in each login, he will —instead— wait, say 10 seconds in his initial login in order to utilise a much more expensive key derivation, and then wait near 0 seconds for every subsequent login as the expensive key is cached.

- derived keys can be cached securely, without increasing most users' assumptions. e.g. cached keys can live in a *dm-crypt* partition that is encrypted with a large encryption key that is stored properly, and the cache can have strict read permissions so that only the unique user that runs **ciphart** executable can read it.

this only requires to trust the user **root**, which is already trusted by almost everyone. so we are not introducing a new assumption.

¹⁰unless hashing collisions happen, which is very unlikely.

for most people, if `root` is compromised, then the adversary can break every other key derivation function, including those that do not cache keys, by simply, say, running a keylogger.

so, practically, we are not increasing the assumptions, but we are only increasing the value that we can extract from the assumptions that we already have.

most importantly, utilising discovery 1 allows us to achieve a memory-harder key derivation in an extremely usable way. more details on memory-hardness later.

2.2.2 potential adversary strategies

let's see what may the adversary try to do against a key caching system:

- adversary strategy 1 — hack into user's system and execute a program as that user that tries to read the password cache file to obtain the memory-harder key inside it.

answer: he will not be able to read the file due to strict read permissions of the cached files as set earlier.

- adversary strategy 2 — steal user's hard disk and try to mount it in his system, to login as root and change permissions of the key cache files.

answer: the partition where the cached files are saved are properly encrypted, so he can't see the cached files, let alone changing their read permissions.

- adversary strategy 3 — break into machine's `root` account.

answer: he will succeed, but then he can also run a keylogger, which will also break every other key derivation function, even those who do not use key caching, such as *scrypt*, *argon2*, etc.

2.3 “hashing is only for compression” law

which one is simpler when, say, building a wooden house?

- option 1 — use only nails and, when you need screws, modify some nails into screws.
- option 2 — use nails and screws.

on the surface, option 1 may appear as the simpler choice as it only uses nails, while option 2 uses both nails and screws.

but a deeper look shows that option 1 is actually a lie, as it is also using screws alongside nails, except that the screws are constrained by being re-invented by modifying nails. in other words, option 1 has the extra assumption that its screws must be made using nails, while option 2 does not have this extra assumption. hence option 2 is actually simpler.

my answer with *ciphart* is that option 2 is the simpler choice because it removes the re-invention aspect, specially if the re-invented screws were worse than screws that were made as screws from the start.

on the other hand, when *argon2* acts as if option 1 is better, which is wrong at every level.

before i write about how *argon2* is an example of adopting the mistake in option 1, i want to define the nails and the screws of a key derivation function, strictly from the perspective of key derivation functions.

definition 2 (hashing functions as seen by key derivation functions). *a function that maps input to output such that:*

- *unlimited input* — input is an unbounded number of concatenated data chunks.
- *compression* — input's shannon's entropy bits could be larger than the total bit size of the output.
- *preserving entropy is tried* — output should have as much entropy bits from the input, but entropy loss is possible.
- *walking backwards is extremely hard* — analysing the output to find the input is computationally too hard.

definition 3 (symmetric block encryption functions as seen by key derivation functions). *a function that maps input to output such that:*

- *limited input* — fixed sized data and a fixed sized key.
- *preserving entropy is guaranteed* — no entropy loss is possible. the proof is that, if the encryption key is known, we can bring back every input bit from the output by decryption.
- *walking backwards is extremely hard* — analysing the output to find the input is computationally too hard.

when *argon2* completes solving the tasks in its memory pad, it derives the output key by hashing certain chunks of bytes in the pad. the number of hashed chunks depends on pad's size, so it's not of a fixed size, and henceforth meets the properties of a hashing function shown in definition 2. therefore, *argon2* using a hashing function at this stage is justified.

when *argon2* is solving tasks in the memory pad, it still uses a hashing function, despite the fact that it is strictly dealing with inputs of fixed lengths, which also meets the properties of a symmetric block encryption function in definition 3. here, *argon2* better use a symmetric block encryption function instead of a hashing function. using a hashing function here is a problem due to:

- re-invention of wheels — having *argon2* concatenate two inputs of a fixed size in order to derive an output of the same size is effectively an attempt to re-invent a symmetric block encryption function. i.e. the concatenation is used to emulate the effect of having a pre-shared key which exists in symmetric block encryption functions. why re-invent keys by concatenation when there exists functions that already have keys?

- needless risk of entropy loss — using a hashing function when dealing with fixed input sizes needlessly increases the probability of having potential entropy losses. this is due to the fact that hashing functions *try* to preserve input’s entropy, but cannot guarantee it, while symmetric block ciphers *do* guarantee it. so why have the possibility of losing entropy bits when you don’t have to?

- slower memory filling rate — generally speaking, hashing functions tend to be slower than symmetric block encryption functions. this is because of dealing with compression is harder than not.

symmetric block ciphers guarantee preserving input’s entropy in the output with much less effort thanks to the fact that the output is at least as large as the input.

but hashing functions don’t have the luxury of having an output that’s as large as the input, thus they need to work a lot more in order to ensure that no input entropy is needlessly lost.

this slowness is bad as it reduces number of passes over the memory pad in a unit of time. more passes over the memory pad are important for strengthening the memory hardness.

law 1 (hashing is only for compression). *use hashing functions only when compression happens. otherwise, use symmetric block encryption functions.*

3 ciphart

3.1 the algorithm

3.1.1 parameters

P	password.
S	salt.
M	total random-access memory in bytes.
D	total hard-disk memory in bytes.
F	temporary file’s path.
Y	whether key caching is enabled.
L	number of memory lanes for concurrency.
T	number of tasks per lane segment.
H	number of lane segments per hard-disk read.
B	minimum <i>caveman’s entropy bits</i> to inject into p .
K	output key’s size in bytes.

3.1.2 internal variables

enc	encryption function.
hash	hashing function.
read	hard-disk file reading function, with seeking. e.g. read (x, y, z) reads z many bytes from file x after seeking y bytes forward.
write	hard-disk file writing function.
C	$\leftarrow \begin{cases} 64 \text{ bytes} & \text{if } \mathbf{enc} \text{ is } xchacha20 \\ 16 \text{ bytes} & \text{if } \mathbf{enc} \text{ is } aes \\ \vdots \end{cases}$ this to reflect the block size of the encryption algorithm that implements enc .
V	$\leftarrow \begin{cases} 32 \text{ bytes} & \text{if } \mathbf{enc} \text{ is } xchacha20 \\ 16 \text{ bytes} & \text{if } \mathbf{enc} \text{ is } aes-128 \\ 32 \text{ bytes} & \text{if } \mathbf{enc} \text{ is } aes-256 \\ \vdots \end{cases}$ this is the size of the encryption key that’s used to solve <i>ciphart</i> ’s tasks. this is different than output key’s size, K , which is enc -independent.
\hat{T}	$\leftarrow \max(\lceil VC^{-1} \rceil, T)$. this is to ensure that we have enough encrypted bytes for new keys.
\hat{T}	$\leftarrow \hat{T} - (\hat{T} \bmod 2) + 2$. this is to ensure that there is an even number of tasks in a segment. why? because we need a buffer for storing the clear-text and another for storing the output cipher-text.
\hat{M}	$\leftarrow M - (M \bmod C\hat{T}L) + C\hat{T}L$. this is to ensure that it is in multiples of $C\hat{T}L$. why? so that all segments are of equal lengths in order to simplify <i>ciphart</i> ’s logic. e.g. it wouldn’t be nice if the last segments were of unequal sizes.
G	$\leftarrow \hat{M}C^{-1}\hat{T}^{-1}L^{-1}$. total number of segments per lane.
n	$\leftarrow 0$. actual number of times enc is called.
m_i	C -bytes memory for i^{th} task in the \hat{M} -bytes pad.
n_l	$\leftarrow \max(\text{nonce})L^{-1}l$. nonce variable for l^{th} lane.
f	$\leftarrow 0$. a counter indicating number of times memory is filled with \hat{M} many bytes.
d	$\leftarrow 0$. a counter indicating total number of saved blocks into hard-disk.
h	$\leftarrow 0$. a counter indicating number of processed lane segments since the last hard-disk read.
u	$\leftarrow 0$. a counter indicating number of times key was updated from the hard-disk.
$*v$	$\leftarrow *hash(P\ S\ M\ D\ \dots\ K, V)$. a pointer to the first byte where V -bytes key is stored. v is the key itself, and $*v$ is a pointer to it.
Z	$\leftarrow hash(v\ 0, V)$. file name where output key, k , is expected to be cached, if k was previously cached.

3.1.3 output

- k K -bytes key.
- n total number of times **enc** was actually called.
 $\log_2 n$ is total number shannon's entropy bits that *ciphart* injected into k , such that $\log_2 n \geq B$.
- \hat{M} actual number of bytes required to exist in random-access memory.
- d actual number of bytes required to exist in hard-disk.

3.1.4 steps

steps of *ciphart* is shown in algorithm 1. this corresponds to *argon2d*. adding a *ciphart-i* variant is a trivial matter, i just didn't do it yet because my threat model currently doesn't benefit from a password independent variant.

3.2 noteworthy features

3.2.1 parallelism

since iterations of the loop in line 6 in algorithm 1 are fully independent of one other, they can quite happily utilise L cpu cores, specially when segment sizes, T , are larger.

3.2.2 memory-hardness

Proof. algorithm 1 is just a variation of *argon2d*. so if *argon2d* is memory-hard, then so is *ciphart*. \square

3.2.3 memory-hardness

thanks to discovery 1, we can cache keys, as done in line 1, without increasing assumptions of the threat model of the vast majority of users. then, memory-hardness becomes possible. this process goes like this:

1. starts by running an **enc**-based variant of *argon2*, except that, as it is going, it keeps writing the updated segments into the hard-disk until the size of it satisfies the D bytes limit. this is shown in line 26 in algorithm 1.

optimising this hard-disk filling with D bytes is not a big deal, since this feature is probably going to be used only once; thanks to key caching. that said:

- this feature is optional. i.e. in case someone doesn't like the hard-disk caching, he can set $D \leftarrow 0$ to disable it. but, for most people, i don't understand why you would want to disable it. e.g. if you're already trusting **root**, then i think that you can use this feature without changing your threat model.

i personally like it a lot as it allows me to achieve memory-hardness way beyond my random-access memory. just imagine the look on the face of those asics crackers once they hear that your *ciphart* requires, say, 50 giga bytes!

algorithm 1: ciphart

```

1 if  $Y$  and exists( $Z$ ) then
2    $k, n, d \leftarrow \text{read}(Z, 0, V + \text{sizeof}(k \| n \| d))$ ;
3   return  $k, n, d$ 
4 while 1 do
5   for  $g \leftarrow 0, 1, \dots, G - 1$  do
6     for  $l \leftarrow 0, 1, \dots, L - 1$  do
7       for  $t \leftarrow 0, 1, \dots, T - 1$  do
8          $i \leftarrow gLT + lT + t$ ;
9         if  $t < T - 1$  then
10           $j \leftarrow i + 1$ ;
11        else if  $t = T - 1$  then
12           $j \leftarrow i - T + 1$ ;
13         $m_j \leftarrow \text{enc}(m_i, n_l, *v)$ ;
14         $n_l \leftarrow n_l + 1$ ;
15        if  $f = 0$  then
16           $*v \leftarrow m_j \bmod (gLTC + tC - V)$ ;
17          if  $*v \geq gLTC - V$  then
18             $*v \leftarrow *v + LTC$ ;
19          else
20             $*v \leftarrow m_j \bmod (\hat{M} - LTC + tC - V)$ ;
21            if  $*v \geq gLTC + tC - V$  then
22               $*v \leftarrow *v + LTC$ ;
23            else if  $*v \geq gLTC - V$  then
24               $*v \leftarrow *v + LTC$ ;
25         $n \leftarrow n + LT$ ;
26        if  $d \leq D$  then
27          for  $i \leftarrow gLT, \dots, gLT + (T - 1)$  do
28            write( $F, m_i$ );
29             $d \leftarrow d + 1$ ;
30            if  $d \geq D$  then
31              break;
32        else
33           $h \leftarrow h + L$ ;
34          if  $h \geq H$  then
35             $*v \leftarrow * \text{read}(F, v \bmod (d - V), V)$ ;
36             $u \leftarrow u + 1$ ;
37             $h \leftarrow 0$ ;
38          if  $f \geq 1$  and  $u \geq 1$  and  $n \geq 2^B$  then
39             $g_{\text{last}} \leftarrow g$ ;
40            go to line 42;
41         $f \leftarrow f + 1$ ;
42     $i \leftarrow g_{\text{last}}LT$ ;
43     $k \leftarrow \text{hash}(m_{i+0T} \| m_{i+1T} \| \dots \| m_{i+(L-1)T}, K)$ ;
44    if  $Y$  then
45      write( $Z, k \| n \| d$ );
46    delete( $F$ );
47    return  $k, n, d$ 

```

- when key caching is enabled, i.e. $Y \leftarrow 1$, this hard-disk writing is done only initially, and subsequent uses appear as almost instantaneous.
- this hard-disk writing can be slightly optimised by using a non-blocking write operation. but i think that the trivial reduction in this time may not justify increasing code's complexity, so i don't plan to implement non-blocking writes in `libciphart`.

2. then, once the D bytes hard-disk requirement is satisfied, the process continues to run the modified `enc`-based variant of `argon2` except that it updates the key v from those D bytes every time S many segments are solved. this step makes `ciphart` require $D + M$ bytes. this is shown in line 34 in algorithm 1.

if S is large enough, this can be done efficiently without blocking the cpu noticeably, as the randomly obtained V bytes can be read by using the $O(1)$ operation, `seek`, over the D bytes.

3. delete the D bytes, just to free up the disk space. no need to securely delete those D bytes since we can save them using a temporary random key that's forgotten later on.

3.3 comparison

4 application scenarios

4.1 a currently-useful scenario

user has a password manager which generates unique 256 bit entropy keys for each online service that he uses. the user also renews keys for his online services every now and then. so his online accounts generally have high security.

but user's problem is how to lock and unlock his password manager's passwords database. should he use a physical usb-stick key that types a high-entropy key that encrypts and decrypts the database? the user doesn't want this physical key because of several reasons:

- he tends to lose his keys a lot, and, for certain tasks, the risk of needing to wait for until he gets a backup usb-stick key is too much.
- he doesn't want to be caught having cryptographic usb-stick keys, because as such is an evidence that he has encrypted content. the user wants to have the choice to lie that he has no clue. so not having usb-stick keys with him helps his case to lie.
- he wants to be torture-resistant so that an adversary cannot forcefully take his keys from him in order to login into his services. he may rather want to die than to give the password to the adversary. this works because, so far, the brain is a pretty private information store.

in this scenario, the user memorises a sensible password that he can remember, with enough initial entropy, and then uses `ciphart`, preferably with disk caching, to inject large amounts of entropy bits into his derived keys, way beyond the reach of pre-`ciphart` key derivation functions; thanks to theorem 1 and discovery 1.

here, the expensively derived key is only used to unlock a local password manager, which offers a protection against situations where a backup copy of the passwords database is stolen. this may enable the user to store his passwords manager in an online file synchronisation service for more convenient system migrations to further reduce his login delays should he face the need to migrate to a new, say, personal computer.

advertisement 1. *in case you're interested in such a passwords manager, i've also made nsapass¹¹ — a flexible and a simple passwords manager in a few hundreds of python lines of code, that uses ciphart by default. i think this is the best command-line interfacing passwords manager by far, for its usability, and for the fact that auditing it is easy, thanks for it being only in a few hundreds of python lines.*

4.2 a later-useful scenario

all input password fields will internally call `libciphart` to derive more expensive keys. this way, applications, such as `firefox`, `mutt`, ..., will never send actual passwords, but will only send `ciphart`-derived keys with increased shannon's entropy content.

thanks to theorem 1, this will automatically increase shannon's entropy of all users' passwords without requiring users to memorise harder passwords. thanks to discovery 1, the user will not face any delay in his daily use, except only an initial delay to create the cache entry.

i also think that it is *generally* better to have expensive key derivation functions in the client side as opposed to the server side. because remote servers always have the incentive to reduce the complexity of the key derivation function in order to free more resources for other things that bring them money.

either way, `ciphart` is perfectly usable on the server side. it is just that i think `ciphart`, `argon2`, `scrypt`, ... make better sense when placed on the client side.

¹¹<https://github.com/Al-Caveman/nsapass>

A donations

this work is sponsored by an unexpected nerdiness. nothing here was supposed to happen. the cause remains unknown, so we call it *something random with a large entropy*.

ancient mythology has it that, every time a donation is made for a good cause, a beautiful torch is kindled somewhere deep within an otherwise cold and dark space.

the torch is pure with an innocent smile every time a spending is committed towards good. but, she drops a tear when the spending is done in excess. despite her attempts to maintain a steady light, sometimes her tears force her light to flicker, and sometimes the flickering causes her to disappear. when she is gone, she is never back again, except in the memories of those who have once witnessed her charm...

so, if you opt for a donation, i'd appreciate exercising wisdom, so that the sensible balance is not exceeded; lest we want a beautiful torch to be born, only to make her disappear in her tears.

A.1 bitcoin



bc1qtylztgd0yu4v7f8hyfzufn7nu692v9fc88jln