ciphart

memory-hard key derivation
with easier measurable security

caveman[1]
2021-01-10 12:07:05+00:00

---

argon2[2] is mostly nice, but trying to interpret its contribution to the protection against password brute-forcing attacks remains more difficult than it should be. this vagueness is a problem that is not limited to *argon2*, but also shared with every other key derivation function that i've known so far.

when one uses *argon2*, his derived key will surely have superior protection against password brute-forcing attacks, but by how much? to answer this, one would need to survey the industry that manufactures application-specific integrated circuits to obtain a map between *time* and *money*, in order to get an estimation on how much would it cost the adversary to discover the password in a given time window.

resolving this vagueness is not a mere luxury to have, but a necessity for maximising survival, because it hinders the process of studying the cost-value of memory-hard key derivation functions, which, effectively, increases the risk of having a false sense of security.

so i propose *ciphart* — a memory-hard key derivation function with a security contribution that is measured in a unit that i call *relative entropy bits*. this unit is measured objectively and is guaranteed to be true irrespective of whatever alien technology that the adversary might have.

`libciphart`[3] is a library that implements *ciphart* very closely to this paper, without much fluff. this should make integrating *ciphart* into other systems more convenient.

`ciphart`[4] is an application for encrypting and decrypting files that makes use of `libciphart`. this application is intended for use by end-users or scripts, henceforth it has some fluff to treat mankind with dignity.

1 ciphart

1.1 parameters

enc	encryption function.
p	password.
s	salt.
L	number of memory lanes for concurrency.
M	total memory in bytes.
T	number of tasks per lane segment.
R	number of rounds per task.
B	added security in *relative entropy bits*.
K_{out}	output key size in bytes.

---

[1]toraboracaveman [at] protonmail [dot] com
[2]https://github.com/P-H-C/phc-winner-argon2
[3]https://github.com/Al-Caveman/libciphart
[4]https://github.com/Al-Caveman/ciphart

1.2 internal variables

$C \quad \leftarrow \begin{cases} 64 \text{ bytes} & \text{if enc is } xchacha20 \\ 16 \text{ bytes} & \text{if enc is } aes \\ \dots \end{cases}$

this to reflect the block size of the encryption algorithm that's going to use *ciphart*'s generated key to encrypt data.

$K_{in} \quad \leftarrow \begin{cases} 32 \text{ bytes} & \text{if enc is } xchacha20 \\ 16 \text{ bytes} & \text{if enc is } aes\text{-}128 \\ \dots \end{cases}$

this is the size of the encryption key that's used to solve *ciphart*'s tasks. this is different than the enc-independent K_{out} which is possibly used by other encryption algorithms in later stages.

$\hat{T} \quad \leftarrow T - (T \bmod 2) + 2$. this is to ensure that it is in multiples of 2. why? because we need a buffer for storing the clear-text and another for storing the output cipher-text.

$\hat{M} \quad \leftarrow M - (M \bmod C\hat{T}L) + C\hat{T}L$. this is to ensure that it is in multiples of $C\hat{T}L$. why? so that all segments are of equal lengths in order to simplify *ciphart*'s logic. e.g. it wouldn't be nice if the last segments were of unequal sizes.

$G \quad \leftarrow \hat{M}C^{-1}\hat{T}^{-1}L^{-1}$. total number of segments per lane.

$\hat{B} \quad \leftarrow \max(\log_2(\hat{M}C^{-1}R), B)$. this is to ensure that $\hat{B}$ is large enough to have at least one pass over the $\hat{M}$-bytes memory.

$\hat{B} \quad \leftarrow \log_2(2^{\hat{B}} - (2^{\hat{B}} \bmod L\hat{T}R) + L\hat{T}R)$. this is just to reflect the reality with *ciphart* that segments must complete. i.e. when the user asks for B *relative entropy bits*, he gets $\hat{B}$ instead, where $\hat{B} \geq B$. more details on this later.

$m_i \quad C$-bytes memory for i^{th} task in the $\hat{M}$-bytes pad.

$n_l \quad \leftarrow l\hat{T}R$. nonce variable for l^{th} lane with at least 64 bits.

$f \quad \leftarrow 0$. a flag indicating whether the $\hat{M}$-bytes pad is filled.

1.3 output

k $\quad K$-bytes key with $\geq B$ *relative entropy bits*.

1.4 steps

shown in algorithm 1.

2 parallelism

since iterations of the loop in line 3 in algorithm 1 are fully independent of one other, they can quite happily utilise L cpu cores, specially when segment sizes, T, are larger.

other lines are not easily independent, so i didn't even bother to try to parallelise them. specially since this is

---
algorithm 1: ciphart version 6

```
1  while 1 do
2  │   for g = 0, 1, ..., G − 1 do
3  │   │   for l = 0, 1, ..., L − 1 do
4  │   │   │   for t = 0, 1, ..., T − 1 do
5  │   │   │   │   for r = 0, 1, ..., R − 1 do
6  │   │   │   │   │   i ← gT + t;
7  │   │   │   │   │   if t = 0 then
8  │   │   │   │   │   │   j ← i + T − 1;
9  │   │   │   │   │   else if t = T − 1 then
10 │   │   │   │   │   │   j ← i + 1 − T;
11 │   │   │   │   │   else
12 │   │   │   │   │   │   j ← i + 1;
13 │   │   │   │   │   m_j ← enc(m_i, n_l, k);
14 │   │   │   │   │   n_l ← n_l + 1;
15 │   │   │   │   │   k ← f(m_j, p, l, s, t);
16 │   │   if f and log_2(n_1 L) ≥ B then
17 │   │   │   go to line 19;
18 │   f ← 1;
19 while 1 do
20 │   for l = 1, 2, ..., L do
21 │   │   if len(k) ≥ K then return k[0 : K];
22 │   │   n ← n + 1;
23 │   │   k ← k ‖ enc(m_{l,S,T}[1], n, k);
```

not a problem, since the cpu-heavy part is in in the easily-paralleliseable part.

with *argon2*, if one wants to increase the cpu load without increasing memory pad's use, one can increase the number of passes over the pad. this feature is supported by *ciphart* via the *relative entropy bits* parameter B.

but, simply increasing number of passes on the pad may not be the best option for all cases. e.g. what if someone has a small pad, and small segments? in such case, a higher percentage of the cpu will be wasted in the non-paralleliseable steps, which is a waste.

this is why *ciphart* has an additional parameter R. this parameter can allow to increase the load on the cpu without even requiring to go through the non-paralleliseable steps. *argon2* lacks this parameter. this is not the main reason *ciphart* was made, but it's one of the incremental improvements.

one could philosophically argue that *argon2* has the R parameter, except that it always assumes that $R = 1$. i don't agree with this assumption, which is why i made *ciphart* to allows the user to set R more flexibly.