

key derivation with easier measurable security

caveman

January 2, 2021

hi — i propose *ciphart*, a sequential memory-hard key derivation function that has a security gain that's measurable more objectively and more conveniently than anything in class known to date.

to nail this goal, *ciphart*'s security gain is measured in the unit of *relative entropy bits*. relative to what? relative to the encryption algorithm that's used later on. therefore, this *relative entropy bits* measure is guaranteed to be true when the encryption algorithm that's used with *ciphart* is also the same one that's used to encrypt the data afterwards.

my reference implementation is available here¹.

content

1 ciphart

2 parallelism

3 sequential-memory hardness

3.1 example

4 security interpretation

5 comparison

6 summary

1 ciphart

parameters:

M each task's size, at least 32 bytes.
 W total memory in multiples of $2M$.
 R number of rounds per task.
 B added security in *relative entropy bits*.
enc encryption function.
 k initial key.

input:

T $\leftarrow W/M$
 P $\leftarrow \max(2, \lceil 2^B / (TR) \rceil)$
 x $\leftarrow 0$, a 16 bytes wide variable.
 m_t for any task $t \in \{1, 2, \dots, T\}$, m_t is M -bytes memory for t^{th} task to work on. $m_t[0 : 16]$ means first 16 bytes. $m_t[-16 :]$ means last 16 bytes.
nonce a variable with enough bytes to store nonces in.
hash a function to compress W bytes into desired key length.

output:

\hat{k} better key, with B , or more, *relative entropy bits*. specifically, with $\log_2(PTR) \geq B$ bits.

steps:

```
1: for  $p = 1, 2, \dots, P$  do
2:   for  $t = 1, 3, \dots, T - 1$ , in steps of 2 do
3:      $i \leftarrow t$ 
4:      $j \leftarrow t + 1$ 
5:     for  $r = 1, 2, \dots, 2R$  do
6:       nonce  $\leftarrow (p, t, r)$ 
7:        $m_i \leftarrow \text{enc}(m_j, \text{nonce}, k)$ 
8:        $\hat{i} \leftarrow i$ 
9:        $i \leftarrow j$ 
10:       $j \leftarrow \hat{i}$ 
11:    end for
12:     $x \leftarrow x \oplus m_i[-16 :]$ 
13:     $x \leftarrow x \oplus m_j[-16 :]$ 
14:  end for
15:  for  $t = 1, 2, \dots, T$  do
16:     $m_t[0 : 16] \leftarrow m_t[0 : 16] \oplus x$ 
17:  end for
18: end for
19: return  $\hat{k} \leftarrow \text{hash}(m_1, m_2, \dots, m_T)$ 
```

¹<https://github.com/Al-Caveman/ciphart>

2 parallelism

iterations inside the `for` loop, in step 2, are independent of one another, so we can distribute them happily across different threads to achieve maximum cpu utilisation.

exception is in steps 12 and 13, where a mutex might be required as the xor assignment there is not necessarily atomic. but this is not a real problem since the real expensive part is in step 7, which overshadows the overhead of the mutexes.

plus, even if one has an ultra-fancy hardware where step 7 is so lightweight that it effectively competes with the mutexes needed for steps 12 and 13, then one can simply increase number of rounds R until the correct order is restored.

3 sequential-memory hardness

all of the tasks solved in the first pad, i.e. when $p = 1$, can be normally computed sequentially with just, say, $2M$ bytes memory.

sequential-memory hardness is introduced for later pads, $p \geq 2$, as the x variable gets applied to tasks' memory workspaces m_1, m_2, \dots, m_T , which causes every task's memory in pad p to depend on every task's last 16 bytes of the previous pads $p - 1, p - 2, \dots, 1$.

3.1 example

say that we've got a total memory of $W = 10 \times 2M$ bytes, which basically means we have 10 pairs of tasks (or $T = 20$ tasks) in a pad.

also say that we've got B large enough that caused us to need 5 pads, i.e $P = \max(2, \lceil 2^B / (TR) \rceil) = 5$.

for simplicity, let's say $R = 1$. also, say that the adversary has enough memory to keep track of x values across the different pads. since each x is only 16 bytes, and since we've got 5 pads, this means that the adversary has managed 16×5 bytes to not worry about losing x es across the pads.

the question of this example is: how many times will the function `enc` be called if we wanted to complete the ciphart algorithm with just $W = 2M$ bytes for the pad instead of $W + 10 \times 2M$ bytes?

1. in the first pad, $p = 1$, `enc` will be called 20 times. by the end of it, we will x populated with all xor-ed data from all the 20 tasks, and end up having the

content of the two tasks, say m_{19} and m_{20} . but we won't have the content of the other tasks, since we said we have only $2M$ bytes pad size limitation in this example.

2. in the second pad, $p = 2$, we will be able to start working on tasks $t = 19$ and $t = 20$ since we already have their memory content m_{19} and m_{20} as well as the xor-ed data x . so with these two tasks, life is easy. but what about other tasks, say:
 - (a) tasks $t = 1$ and $t = 2$? we already have x , but we lack their memory content from the previous pad. so we've got to repeat that. meaning `enc` will be called 2 times from the previous pad, and then 2 times for the current pad $p = 2$. totalling 4 calls.
 - (b) tasks $t = 3$ and $t = 4$? same as before, 4. this repeats to the rest of task pairs (except the lucky tasks $t = 19$ and $t = 20$).

meaning, we've got $20 - 2$ unlucky tasks, each of which will result in calling `enc` two times in order to solve pad $p = 2$ and obtain its xor-data x .

in total, the second pad will call `enc` function $2 + (20 - 2) \times 2 = 38$ times.

3. in the third pad, $p = 3$, the same will repeat, except for calling `enc` one more time with the unlucky tasks. i.e. $2 + (20 - 2) \times 3 = 56$.
4. in the forth pad, $p = 4$, $2 + (20 - 2) \times 4 = 74$ times.
5. in the fifth pad, $p = 5$, $2 + (20 - 2) \times 5 = 92$ times.

in total, a $W = 2M$ bytes implementation would end up calling `enc` $20 + 38 + 56 + 74 + 92 = 280$ times. **correction:** i think it will be $20 \times 5 + (20 - 2) \times 5 = 190$.

a $W = 10 \times 2M$ bytes implementation will call `enc` $20 \times 5 = 100$ times instead.

so.. is ciphart really hard? it is not. increase in calculation is linear (not exponential). this is a failure. there needs to be an algorithmic change.

4 security interpretation

it sucks. don't use it yet. i'm thinking how to fix this garbage.

5 comparison

6 summary